



deeplearning.ai




Setting up your optimization problem

Normalizing inputs

speed up your learning

Normalizing training sets

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$



Now the variance of x_1 and x_2 are both equal to one.

You just move the training set until it has zero mean.

Subtract mean:

$$\bar{\mu} = \frac{1}{n} \sum_{i=1}^n x^{(i)}$$

$$x := x - \mu$$

Normalize variance

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n x^{(i)} * x^{(i)T}$$

element-wise Squaring

Notice we've already subtracted out the mean, so $x^{(i)}$ squared, element by element is just the variances.

$$x / \sigma^2$$


Use same μ σ^2 to normalize test set.

Because you want your data, both training and test examples, to go through the same transformation defined by the same μ and σ^2 calculated on your training data.


Why normalize inputs?

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$


Unnormalized:



Normalized:



But the rough intuition that your cost function will be more round and easier to optimize when your features are all on similar scales, not from 1 to 1000, zero to one, but mostly from -1 to 1 or of about similar variance of each other. That just makes your cost function J easier and faster to optimize.



$x_1: 0 \dots 1$
 $x_2: -1 \dots 1$
 $x_3: 1 \dots 2$

By just setting all of them to a 0 mean and say, variance 1, that just guarantees that all your features are on a similar scale and will usually help your learning algorithm run faster.

